

Zhe Yu

Research Assistant | Trustworthy Language Models and Agentic AI

HKUST (Guangzhou) | LAI Lab · Zhejiang University | IFRC Lab

zyu@zju-if.com · Homepage · Google Scholar · GitHub

Research Profile

My research focuses on trustworthy language models and agentic AI systems. I study how models internally represent knowledge, evidence provenance, conflict, and risk; why these signals often fail to guide generation and tool use; and how white-box monitoring, audit planning, and verifiable safeguards can improve reliable action and deployment.

Current Research Appointments

The Hong Kong University of Science and Technology (Guangzhou) <i>Research Assistant, LAI Lab; advised by Prof. Chengwei Qin</i>	May 2026 – Present <i>Guangzhou, China</i>
Binjiang Institute of Zhejiang University <i>Research Intern, IFRC Lab; supervised by Dr. Meng Han and Dr. Wenpeng Xing</i>	Nov. 2025 – Present <i>Hangzhou, China</i>
• Research spans trustworthy language models, RAG faithfulness, agentic AI safety, audit reliability, and verifiable deployment under provincial and national key R&D programs.	

Education and Visiting Experience

Communication University of Zhejiang <i>B.Eng. in Artificial Intelligence; undergraduate advisor: Dr. Hao Zeng</i>	2023 – Expected 2027 <i>Hangzhou, China</i>
University of Malaya <i>Visiting Student</i>	Jan. 2025 – Feb. 2025 <i>Kuala Lumpur, Malaysia</i>
Westlake University <i>Visiting Student; supervised by Dr. Ziyang Zhang</i>	Mar. 2024 – Sep. 2024 <i>Hangzhou, China</i>

Selected Research Contributions

- **Retrieval grounding and hallucination detection.** Led problem formulation, method design, experiments, and writing for DISF, a dual-path white-box framework that models Conflict, Drift, and Instability in retrieval-augmented generation; accepted at ACL 2026.
- **Evidence attribution and intervention.** Developed internal-signal methods for retrieval-memory conflict, including FIDES for token-level training-free intervention and the Attribution Blind Spot for distinguishing parametric-memory and retrieved-evidence routes behind identical outputs.
- **Audit and evaluator reliability.** Extended hallucination scores into parent-preserving audit plans with PLAITS and studied why item calibration cannot certify experimental contrasts with GAVEL’s paired audits and pair-level correction.
- **Representation-action dissociation.** Led mechanistic studies showing that source and conflict information can be linearly available before it influences tool choice or answer policy, including Knowing Is Not Acting and Whose Thoughts?; both are under review at NeurIPS 2026.
- **Monitoring and verifiable deployment.** Led LatentAudit and RETINA-SAFE / ECRT for real-time white-box monitoring and high-risk medical triage; contributed to privacy-preserving ownership verification in ZK-FPE, accepted at ACM TURC 2026.

Peer-Reviewed Publications

1. **Zhe Yu***, Wenpeng Xing*, Wenjie Luo, Weize Xu, Lingtong Huang, Yourong Chen, Changting Lin, Meng Han. *DISF: Detecting Hallucinations in Retrieval-Augmented Generation via Dual-path Internal State Forcing Framework*. **ACL 2026**. [PDF]
2. Zhiguo Ma*, **Zhe Yu**, Wenpeng Xing, Yourong Chen, Meng Han. *ZK-FPE: Blockchain-Verifiable Model Fingerprinting with Zero-Knowledge Privacy for Ownership Attribution*. **ACM TURC 2026**. [PDF]
3. Weiping Yu, Weihang Wang, Mingyuan Yan, Keyang He, **Zhe Yu**, Wenpeng Xing, Liyuan Liu, Meng Han. *Trusted Metadata-Coordinated Tiered Off-Chain Storage for Recovery-Safe and Low-Latency IoT Data Management*. *Electronics* (MDPI).
4. F. Zhou, C. Chang, Q. Chang, H. Zhang, **Zhe Yu**, W. Liu, J. Li, J. Yang. *Orthogonal salinity and temperature detection via paralleled dual all-fiber interferometers*. *Optics Communications*, 583 (2025): 131688. [DOI]
5. **Zhe Yu**, H. Zeng, Y. Zhao, X. Zhang, Z. Wang, Y. Tao, M. Yuan, X. Sun. *Bibliometric analysis of physical education research in China from 2014 to 2024*. *ACM ICETM*, 2025, pp. 128–132. [DOI]

Manuscripts Under Review

1. **Zhe Yu**, Yunzhao Wei, Wenpeng Xing, Wenjie Luo, Zaobo He, Quan Chen, Meng Han. *PLAIT: From Hallucination Scores to Parent-Preserving Audit Plans*. Under review at **AAAI 2027**. [PDF]
2. **Zhe Yu***, Wenpeng Xing*, Bo Yang, Chen Ye, Gaolei Li, Yunzhao Wei, Meng Han. *The Attribution Blind Spot: Detecting When Language Models Rely on Memory Rather Than Retrieved Context*. Under review at **AAAI 2027**. [PDF]
3. **Zhe Yu***, Wenpeng Xing*, Chen Ye, Xuyang Teng, Bo Yang, Changting Lin, Meng Han. *Calibration Does Not Certify a Contrast: Identification and Correction for LLM Evaluators*. Under review at **AAAI 2027**. [PDF]
4. **Zhe Yu**, Wenpeng Xing, Zhenhua Xu, Xingxing Yang, Meng Han. *Knowing Is Not Acting: Representation-Action Dissociation in Indirect Prompt Injection*. Under review at **NeurIPS 2026**. [PDF]
5. **Zhe Yu**, Wenpeng Xing, Zhenhua Xu, Ruiqi Zhang, Meng Han. *Whose Thoughts? Chain-of-Thought Override in Reasoning-Tuned Language Models*. Under review at **NeurIPS 2026**. [PDF]
6. **Zhe Yu***, Wenpeng Xing*, Tiancheng Zhao, Mohan Li, Changting Lin, Meng Han. *FIDES: Faithful Inference via Deep Evidence Signals for Retrieval-Memory Conflict in RAG*. Under review at ARR / EMNLP. [PDF]
7. **Zhe Yu***, Wenpeng Xing*, Gaolei Li, Shuguang Xiong, Hongzhi Wang, Xuyang Teng, Meng Han. *Cordon-MAS: Defending RAG against Knowledge Poisoning via Information-Flow Control*. Under review at ARR / EMNLP. [arXiv]
8. **Zhe Yu***, Wenpeng Xing*, Yunzhao Wei, Hongzhi Wang, Xuyang Teng, Meng Han. *Composition Collapse: Stable Factual Knowledge Does Not Imply Compositional Reasoning*. Under review at ARR / EMNLP. [arXiv]
9. Zhenhua Xu, Qichen Liu, Zhebo Wang, **Zhe Yu**, Xixiang Zhao, Wenpeng Xing, Dezhang Kong, Mohan Li, Meng Han. *Fingerprint Vector: Enabling Scalable and Efficient Model Fingerprint Transfer via Vector Addition*. Under review at ARR / EMNLP.
10. **Zhe Yu***, Wenpeng Xing*, Meng Han. *LatentAudit: Real-Time White-Box Faithfulness Monitoring for Retrieval-Augmented Generation with Verifiable Deployment*. Under review at CoLM 2026. [arXiv]
11. **Zhe Yu***, Wenpeng Xing*, Meng Han. *From Retinal Evidence to Safe Decisions: RETINA-SAFE and ECRT for Hallucination Risk Triage in Medical LLMs*. Under review at MICCAI 2026. [arXiv]

* Equal contribution. Review statuses are a snapshot as of August 2026.

Patents

1. **Zhe Yu**, Wenpeng Xing, Meng Han. *A hallucination detection method based on dual-path internal state forcing logic for retrieval-augmented generation in large language models*. Chinese Patent Application No. 202610260408X. Under review.
2. Meng Han, **Zhe Yu**, Jiayan Hu, Rongchang Li, Wenpeng Xing, Jingyi Yu, Zhen Hong, et al. *A post-processing method, system, device, and medium for hallucination detection in large language models based on adaptive order statistics aggregation*. Chinese Patent Application No. 2026107898102, filed June 3, 2026. Pending.
3. **Zhe Yu**, Jiayan Hu, Jingyi Yu, Weihang Yu, Wenpeng Xing, Jing Xiong, Yourong Chen, Zhen Hong, et al. *A hallucination detection method, system, and device for large language models based on multi-dimensional heterogeneous feature fusion*. Chinese Patent Application No. 2026107899270, filed June 3, 2026. Pending.

Industry Experience

BoostEngine

R&D Intern, Full-Stack Development / AI Agents

Jul. 2025 – Oct. 2025

Hangzhou, China

- Built a multi-brand Manifest V3 Chrome extension for TikTok Shop operations across 20 regional domains and 9 language packs using React, TypeScript, Vite, and Chrome APIs.
- Developed Spring Boot analytics services for 8,000+ creator collaborations and built a FastAPI / Redis / MySQL synchronization system with anti-loop control, queue coalescing, monitoring, and alerts.

Awards and Technical Skills

- **Silver Award**, CCB Cup Zhejiang Provincial International College Students' Innovation Competition, 2024, for the project *Scoliosis Detection AI System*.
- **Research**: Python, PyTorch, Hugging Face Transformers, scikit-learn, NumPy, pandas, FAISS, vLLM, W&B, MATLAB.
- **Engineering**: TypeScript / JavaScript, Java, React, FastAPI, Spring Boot, SQLAlchemy, MyBatis-Plus, MySQL, Redis, Docker, Chrome Extensions.